

# WHO dataset advanced visualization

Luke Geel

11/3/2020

```
library(tidyverse)


## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.4.0      v purrr  1.0.0
## v tibble  3.1.8      v dplyr  1.0.10
## v tidyr   1.2.1      v stringr 1.5.0
## v readr   2.1.3      v forcats 0.5.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()

gss_cat

## # A tibble: 21,483 x 9
##   year marital      age race rincome partyid relig denom tvhours
##   <int> <fct>      <int> <fct> <fct>   <fct>   <fct> <fct>   <int>
## 1 2000 Never married 26 White $8000 to 9999 Ind,near ~ Prot~ Sout~    12
## 2 2000 Divorced     48 White $8000 to 9999 Not str r~ Prot~ Bapt~    NA
## 3 2000 Widowed      67 White Not applicable Independe~ Prot~ No d~     2
## 4 2000 Never married 39 White Not applicable Ind,near ~ Orth~ Not ~     4
## 5 2000 Divorced     25 White Not applicable Not str d~ None  Not ~     1
## 6 2000 Married      25 White $20000 - 24999 Strong de~ Prot~ Sout~    NA
## 7 2000 Never married 36 White $25000 or more Not str r~ Chri~ Not ~     3
## 8 2000 Divorced     44 White $7000 to 7999 Ind,near ~ Prot~ Luth~    NA
## 9 2000 Married      44 White $25000 or more Not str d~ Prot~ Other     0
## 10 2000 Married     47 White $25000 or more Strong re~ Prot~ Sout~     3
## # ... with 21,473 more rows
```

Question 1 General Social Survey A. Make a bar chart for the rincome (reported income). What makes the default bar chart hard to understand?

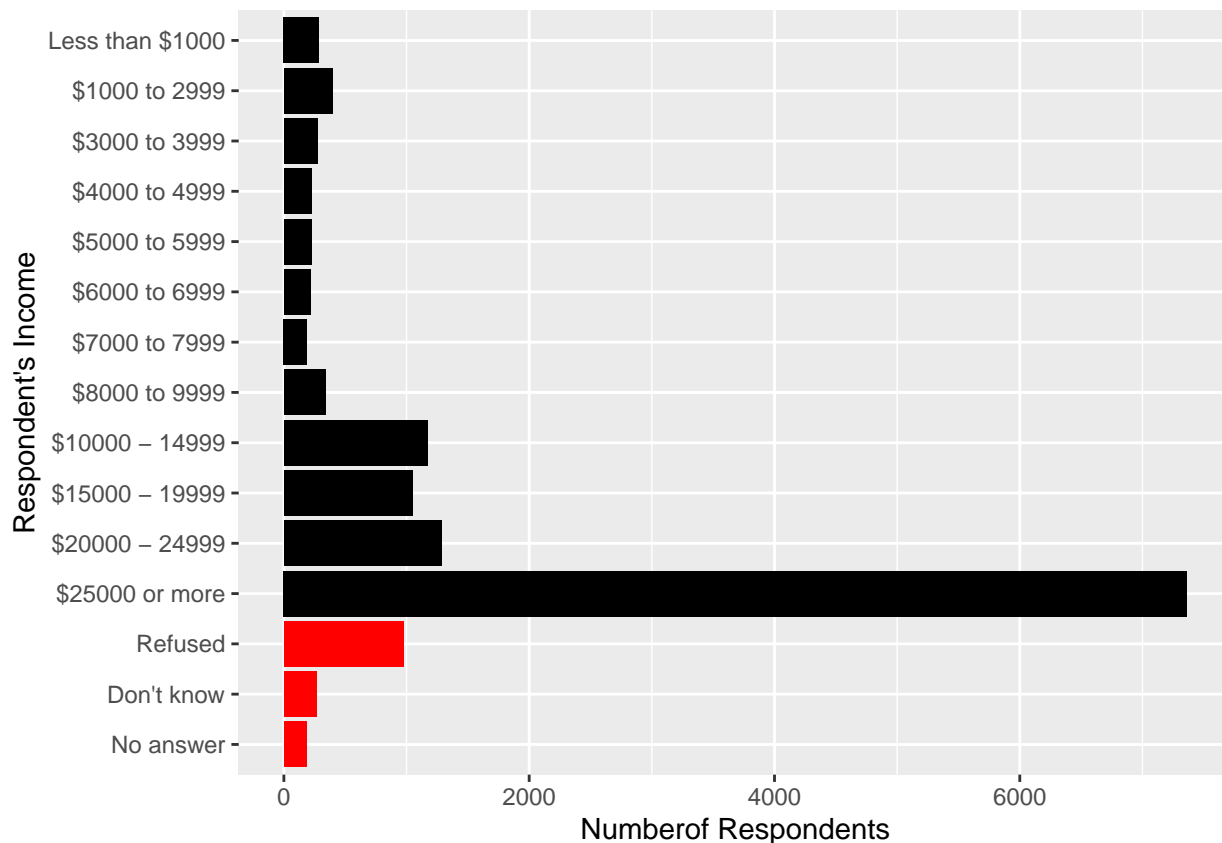
```
ggplot(gss_cat)+
  geom_bar(aes(x=rincome))
```



The reason this default bar chart is so hard to understand is because the x-axis labels are so close together that you can't tell them apart.

B. Let's make this bar chart better: 1. remove the rows with value "Not applicable" 2. rename "Lt \$1000" to "Less than \$1000" 3. (optional) use color to distinguish non-response categories ("Refused", "Don't know", and "No answer") from income levels ("Lt \$1000", ...) 4. add meaningful y- and x-axis titles 5. flip the coordinate

```
new_gss_cat <- gss_cat  
levels(new_gss_cat$rincome)[levels(new_gss_cat$rincome)=="Lt $1000"] <- "Less than $1000"  
ggplot(new_gss_cat %>% filter(rincome != "Not applicable"))+  
  geom_bar(aes(y=rincome, fill=rincome))+  
  labs(x= "Numberof Respondents",y="Respondent's Income")+  
  scale_fill_manual(values = c("red", "red","red","black","black","black","black","black","black","black"),  
    theme(legend.position = "none")
```



Question 2 who Data

```
names(who) <- str_replace(names(who), "newrel", "new_rel")
who2 <- who %>%
  gather("codes", "case", 5:60) %>%
  select(-iso2, -iso3) %>%
  separate(codes, c("new", "type", "sexage"), sep = "_") %>%
  select(-new) %>%
  separate(sexage, into = c("sex", "age"), sep = 1, convert=TRUE)
```

A. There are many missing values in the case variable. We need to think about how missing values are represented in this dataset. The main concern is whether a missing value means that there were no cases of Tuberculosis (TB) or whether it means that the WHO does not have data on the number of TB cases. Check the presence of zeroes in the case variable.

```
nrow(who2%>%filter(case=="0"))
```

```
## [1] 11080
```

```
nrow(who2%>%filter(is.na(case)))
```

```
## [1] 329394
```

Based on these results I believe that a value of 0 means that there were 0 cases while an NA value means that the data is missing.


B. For each year and sex compute the total number of cases of TB (remove the missing values), and make the following time series plot of the total number of cases (Note the changes in the legend labels and the scientific notation in the y-axis). What do you find from the plot?

```

who2 %>%
  filter(!is.na(case)) %>%
  group_by(year, sex) %>%
  summarize(total_cases = sum(case))%>%
  ggplot()+
  geom_line(mapping=aes(x=year,y=total_cases, color=sex))+
  labs(y="total")

```

## `summarise()` has grouped output by 'year'. You can override using the  
## `.groups` argument.



C. Let's plot the number of cases for each age group for males and females.

```

who3 <- who2%>%filter(!is.na(case))%>%group_by(year,age)
attach(who3)
who3$agecat[age == 14 ] <- "0-14"

```

## Warning: Unknown or uninitialised column: `agecat`.

```

who3$agecat[age == 1524 ] <- "15-24"
who3$agecat[age == 2534 ] <- "25-34"
who3$agecat[age == 3544 ] <- "35-44"
who3$agecat[age == 4554 ] <- "45-54"
who3$agecat[age == 5564 ] <- "55-64"
who3$agecat[age == 65 ] <- "65-"

```


```

who3 %>%group_by(year, agecat,sex)%>%
  summarize(total_cases = sum(case),sex=sex)%>%

```

```
ggplot()+
  geom_line(mapping=aes(x=year, y=total_cases, color=agecat))+
  facet_grid(.~sex)+
  labs(y="total", color="age")
```

## `summarise()` has grouped output by 'year', 'agecat', 'sex'. You can override  
## using the `.groups` argument.



D. Collapse the age groups into five categories 0-14,15-24, 25-44, 45-64,65-. Below is an example.

```
who3 <- who2%>%filter(!is.na(case))%>%group_by(year,age)
attach(who3)
```

## The following objects are masked from who3 (pos = 3):

##

## age, case, country, sex, type, year

```
who3$agecat[age == 14 ] <- "0-14"
```

## Warning: Unknown or uninitialised column: `agecat`.

```
who3$agecat[age == 1524 ] <- "15-24"
```

```
who3$agecat[age == 2534 | age == 3544 ] <- "25-44"
```

```
who3$agecat[age == 4554 | age == 5564] <- "45-64"
```


```
who3$agecat[age == 65 ] <- "65-"
```

```
who3 %>%group_by(year, agecat,sex)%>%
  summarize(total_cases = sum(case),sex=sex)%>%
  ggplot()+
```

```
geom_line(mapping=aes(x=year, y=total_cases, color=agecat))+
facet_grid(.~sex)+
labs(y="total", color="age")+
xlim(1990,2013)
```

## `summarise()` has grouped output by 'year', 'agecat', 'sex'. You can override  
## using the `.groups` argument.


## Warning: Removed 133 rows containing missing values (`geom\_line()`).



Question 3: More and more on flights dataset A. Explain what the following code does.

```
library(nycflights13)

airports %>%
  semi_join(flights, c("faa" = "dest")) %>%
  ggplot(aes(lon, lat))+
  geom_point()
```



What this code does is it joins the airports dataset with the flights dataset where the faa airport code from airports matches the dest variable from flights and plots each airport in both datasets.

B. Here's an easy way to draw a map of the United States. (You need to install and load the maps package.) `airports %>% semi_join(flights, c("faa" = "dest")) %>% ggplot(aes(lon, lat)) + borders("state") + geom_point() + coord_quickmap()`

Modify the code to show the spatial distribution of delays. You need to 1. compute the average arrival delay by destination from the flights data 2. join on the airports data to get the longitude and latitude of each destination 3. plot each destination with the average arrival delay mapping to the size or color aesthetic (map a continuous variable to color or size will generate a bar in the legend) 4. add the borders of the states and change the coordinate using `coord_quickmap()`

```
library(maps)


##
## Attaching package: 'maps'
## The following object is masked from 'package:purrr':
##
##      map

#library(plyr)
#library(dplyr)
library(nycflights13)

flights1 <- flights %>% group_by(dest) %>%
  summarize(avg_arr_delay = mean(arr_delay, na.rm=TRUE))

flights1 %>%
```

```
inner_join(airports, c("dest" = "faa")) %>%
ggplot(aes(lon, lat)) +
  borders("state") +
  geom_point(aes(color = avg_arr_delay)) +
  coord_quickmap()+
  labs(color="delay")
```



C.

We need to first get the longitude and latitude of the origin and destination of each flight. You need to 1. Join flights on airports by origin to get the longitude and latitude of the origin 2. Join flights on airports by destination to get the longitude and latitude of the origin

```
airports%>%semi_join(flights, c("faa"="origin"))
```

```
## # A tibble: 3 x 8
```

##	faa	name	lat	lon	alt	tz	dst	tzone
##	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>	<chr>
## 1	EWR	Newark Liberty Intl	40.7	-74.2	18	-5	A	America/New_York
## 2	JFK	John F Kennedy Intl	40.6	-73.8	13	-5	A	America/New_York
## 3	LGA	La Guardia	40.8	-73.9	22	-5	A	America/New_York

```
airports%>%semi_join(flights, c("faa"="dest"))
```

```
## # A tibble: 101 x 8
```

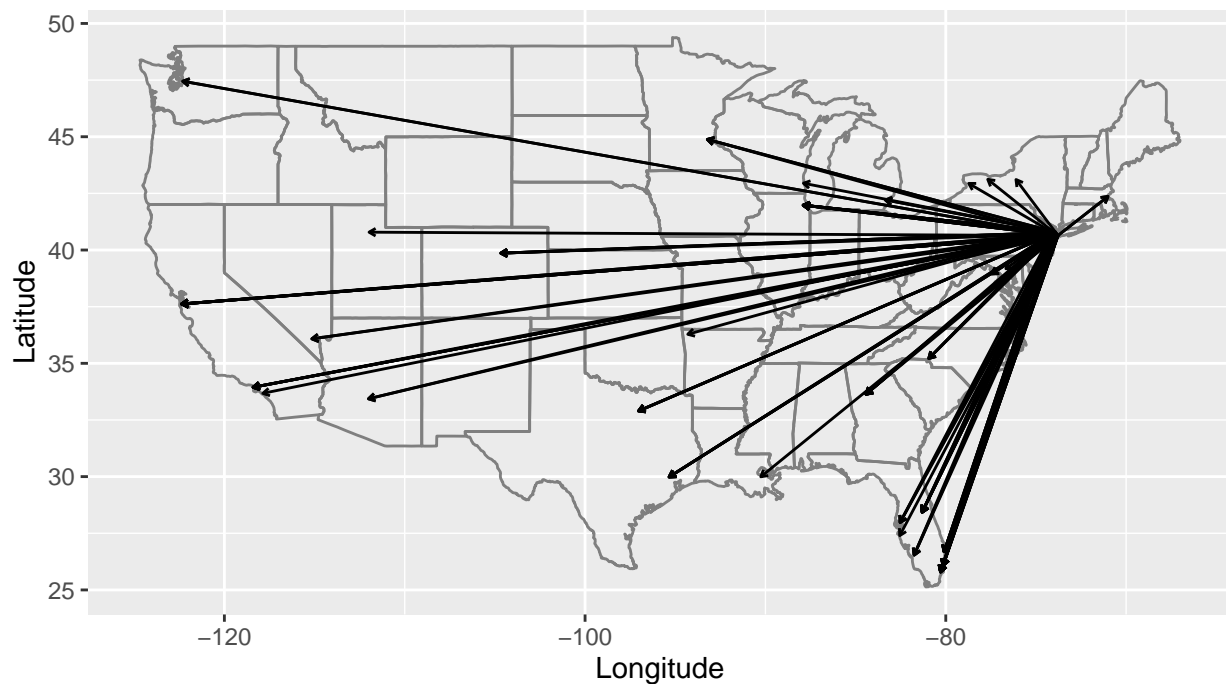
##	faa	name	lat	lon	alt	tz	dst	tzone
##	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>	<chr>
## 1	ABQ	Albuquerque International Sunport	35.0	-107.	5355	-7	A	Ameri~
## 2	ACK	Nantucket Mem	41.3	-70.1	48	-5	A	Ameri~
## 3	ALB	Albany Intl	42.7	-73.8	285	-5	A	Ameri~
## 4	ANC	Ted Stevens Anchorage Intl	61.2	-150.	152	-9	A	Ameri~
## 5	ATL	Hartsfield Jackson Atlanta Intl	33.6	-84.4	1026	-5	A	Ameri~
## 6	AUS	Austin Bergstrom Intl	30.2	-97.7	542	-6	A	Ameri~
## 7	AVL	Asheville Regional Airport	35.4	-82.5	2165	-5	A	Ameri~
## 8	BDL	Bradley Intl	41.9	-72.7	173	-5	A	Ameri~
## 9	BGR	Bangor Intl	44.8	-68.8	192	-5	A	Ameri~
## 10	BHM	Birmingham Intl	33.6	-86.8	644	-6	A	Ameri~



```
## # ... with 91 more rows
```

D. We draw the route only for the first 100 flights: 1. use `slice(1:100)` to extract the first 100 rows from the data obtained in (c) 2. use `geom_segment` to draw the route for each flight. Check `arrow` argument to adjust the size of the arrow 3. add the borders of the states and change the coordinate using `coord_quickmap()`

```
flights %>%
  inner_join(select(airports, origin = faa, origin_lat = lat, origin_lon = lon), by = "origin") %>%
  inner_join(select(airports, dest = faa, dest_lat = lat, dest_lon = lon), by = "dest") %>%
  slice(1:100) %>%
  ggplot(aes(lon, lat)) +
  borders("state") +
  geom_segment(mapping = aes(x=origin_lon,y=origin_lat,xend=dest_lon,yend=dest_lat), arrow = arrow(leng
  coord_quickmap()+
  labs(x="Longitude", y="Latitude")
```



E. Let's explore the relationship between the age of a plane (each plane is represented by `tailnum`) and its delay. The `year` variable in the `planes` data shows the year manufactured. Following the steps to draw a figure (below is an example): 1. join flights on planes (note the flights also has a variable `year`) 2. calculate the age of the plane 3. calculate the average departure delay and arrival delay by age 4. gather the average `arr_delay` and `dep_delay` columns into a key:value column pair names `type` and `delay` 5. recode the levels of `type` column to be "Departure" and "Arrival" 6. create the plot

```
#planes%>%
#left_join(flights,by="tailnum") %>%
# summarize(planeage = year.y-year.x)%>%
# summarize(avg_dep_delay = mean(dep_delay), avg_arr_delay = #mean(arr_delay))%>%
# ggplot()+
# geom_line(aes(x=planeage, y=avg_delay))+
# labs(x="age",y="delay")
```

F. Let's explore the relationship between some weather condition and the delay. We will focus on the departure delay and the weather at the origin. Make the following plots and describe the relationship between precipitation (visibility) and the delay.

```
#full_join(flights,weather, by="origin")%>%  
# group_by(origin)%>%  
# mutate(avg_delay = mean(dep_delay), avg_precip = mean(precip))%>%  
# ggplot()+  
# geom_point(aes(x=avg_precip, y=dep_delay))
```